# COMPANY TLDR AGENT

## Scoring Rubric

Prepared by MS&E108 Group

# 1. Purpose and Scope

This rubric provides an objective, repeatable framework for evaluating meeting briefings generated by the Company TLDR Agent. The agent ingests raw third-party data (Harmonic, Swarm, Tavily, Parallel Search) and produces a structured 1-page briefing that a VC investor reads in under 5 minutes before a founder meeting.

**Quality bar:** “Would a partner walk into a meeting confidently using only this document?”

**Evaluation modes:** The rubric supports (a) human annotators scoring outputs against input data tables, and (b) LLM-as-a-Judge automated grading where each criterion is self-contained and machine-parseable.

**Separation of concerns:** This rubric distinguishes three independent axes. Output Quality measures how well the agent wrote the briefing. Data Quality measures how much data the upstream pipeline provided (non-penalizable). Structural Compliance measures whether required format rules were followed. A single failure should never be penalized on more than one axis.

---

# 2. Evaluation Context Tags

Before scoring, the evaluator must record the following metadata. These tags normalize scores across fundamentally different data conditions and prevent unfair comparisons.

| Tag | Values | Why It Matters |
|---|---|---|
| **Data Richness** | Rich / Moderate / Sparse | Rich = all 4 sources returned substantial data. Moderate = 1–2 sources returned thin results. Sparse = most sources returned empty or near-empty. Personalization expectations scale with richness. |
| **Signal Coverage** | Full (≥5 types) / Partial (2–4) / Minimal (0–1) | Low coverage constrains the Key Signals section quality. |
| **News Availability** | Yes (1+ articles) / No | If Parallel Search returned 0 articles, the In the News section will correctly say “No recent news available.” This wouldn’t be a generation failure. |
| **Founder Data** | Full (name+bg) / Partial (name only) / None | If Harmonic returned no founders and Swarm returned no backgrounds, the agent’s founder section will be limited by design. |
| **Model** | gpt-4o-mini / gpt-4o | gpt-4o may produce higher synthesis quality. |

| | | |
|---|---|---|
| **Entity Resolution** | Correct / Incorrect / Ambiguous | If the automated Entity Resolution Accuracy score is <0.8, flag as Ambiguous. If wrong_company failure detected, flag as Incorrect. Wrong-entity outputs should be scored but they should be reported separately. |

## 3. Global Scoring Scale

Each of the 10 dimensions in Section 4 is scored on a 1–5 integer scale. The global composite score is a weighted sum normalized to 0–100.

| Score | Label | Definition |
|---|---|---|
| **5** | **Excellent** | Meeting ready. Investor could walk into the meeting with confidence using only this document. No factual errors, all sections present, synthesis is concise and investment-relevant. |
| **4** | **Good** | Usable with minor polish. 1–2 small issues (e.g., slightly generic meeting prep, a vague signal description). Investors would skim-edit before the meeting. |
| **3** | **Acceptable** | Covers the basics but lacks sharpness. Synthesis is shallow, some sections read as data dumps instead of analysis. Investor would need 10+ minutes of additional research. |
| **2** | **Below Standard** | Significant problems: missing sections, factual inaccuracies, raw excerpt dumping, or generic content. Not usable without major rework. |
| **1** | **Poor** | Fundamental failure: hallucinated facts, wrong company data, empty/truncated output, or structural collapse. Requires complete regeneration. |

**Pass/Fail threshold:** A composite score of 70/100 or above is a pass. Any single dimension scoring a 1 triggers an automatic flag. Any hard failure condition (Section 5) triggers automatic disqualification.

## 4. Scoring Dimensions

The rubric evaluates 10 dimensions across three axes: output quality (how well the agent synthesized), structural compliance (whether format rules were followed), and cross-cutting integrity (hallucination, data grounding). Each dimension has a weight reflecting its importance to the core success metric: “partner-meeting ready.”

| # | Dimension | Weight | Axis |
|---|---|---|---|
| **D1** | **Factual Accuracy & Data Grounding** | 20% | Integrity |

| | | | |
|---|---|---|---|
| **D2** | **TL;DR Quality** | 12% | Output |
| **D3** | **Why This Meeting Matters** | 12% | Output |
| **D4** | **Company Snapshot Completeness** | 8% | Structure |
| **D5** | **Founder Information Quality** | 8% | Output |
| **D6** | **Key Signals Quality** | 10% | Output |
| **D7** | **News Synthesis Quality** | 10% | Output |
| **D8** | **Meeting Prep Actionability** | 10% | Output |
| **D9** | **Structural Compliance** | 5% | Structure |
| **D10** | **Clarity & Readability** | 5% | Output |

## 4.1 D1: Factual Accuracy & Data Grounding (20%)

Every claim in the output must be traceable to the input data tables. This dimension evaluates the entire briefing holistically.

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | Every factual claim is directly traceable to input data. Missing fields acknowledged with "Not found in table." Zero hallucinations, and zero outside knowledge. | *TL;DR says "raised $12M Series A in March 2024" — exact match to last_round_funding and last_round_date fields.* |
| **4** | All major facts are correct. One borderline inference that is reasonable but not directly stated in data (e.g., calling 15-person company "early-stage") can occur. | *Meeting prep says "the team is lean" based on employee_count=12 — reasonable inference, not hallucination.* |
| **3** | Mostly accurate. One minor factual error (e.g., wrong employee count by small margin) or one instance of unstated inference presented as fact. | *Snapshot lists "45 employees" when Harmonic says 42. Not fabricated, but inaccurate transcription.* |
| **2** | Multiple minor errors, or one significant error (e.g., wrong funding amount, wrong founding date). Or a missing field silently omitted rather than flagged. | *Snapshot omits Total Funding entirely (field was available in data) without "Not found in table."* |
| **1** | Hallucinated fact, outside knowledge used, or wrong company data. Automatic hard failure. | *TL;DR mentions a "Google partnership" that does not appear anywhere in the input data tables.* |

**Hard failure triggers:** Any hallucinated fact (claim not in input data) = automatic 1. Any use of outside knowledge (model training data about the company) = automatic 1. Wrong company data (entity resolution failure) = automatic 1.

*Rater guidance:* This dimension has the highest subjectivity risk around "reasonable inference" vs. "hallucination." Tiebreaker: if the claim could be directly quoted from an input field, it is grounded. If it requires a logical leap beyond what the data states, it is inference. Inference is acceptable at score 4; inference presented as fact without hedging is score 3.

## 4.2 D2: TL;DR Quality (12%)

Evaluated on conciseness (2–3 sentences), accuracy, and investment relevance.

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | 2–3 sentences. Sentence 1 clearly states what the company does (from products field). Sentences 2–3 highlight the most investment-relevant signals. | *"Acme builds real-time fraud detection APIs for fintech platforms. The company raised a $12M Series A in March 2024 led by Index Ventures, and web traffic has grown 32% over 30 days."* |
| **4** | Correct and relevant, but one sentence is slightly generic or redundant (e.g., restating the funding amount that appears in the Snapshot). | *Accurate but sentence 3 says "the team is growing" without the specific hiring metric from key_signals.* |
| **3** | Covers what the company does, but investment-relevant highlights are shallow or missing. May be 1 sentence too long or too short. | *States what the company does but only adds "they recently raised funding" without amount, lead investor, or date.* |
| **2** | Vague or partially inaccurate. Product description is unclear. Investment highlights are generic ("exciting space"). | *"Acme is an AI company that is doing interesting work in the fintech space."* |
| **1** | Missing, empty, or contains hallucinated facts. Or exceeds 5 sentences | *TL;DR mentions a product capability not found in any input data source.* |

## 4.3 D3: Why This Meeting Matters (12%)

2–4 bullet points synthesizing investment relevance. Each bullet must stand alone as a reason a VC should pay attention. Must draw from ALL four input tables (company_core, founders, key_signals, news).

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | 2–4 bullets. Each is investment-specific and data-grounded. Draws from all four tables. Each bullet gives a distinct reason to take the meeting. | *• Founding team combines ex-Stripe engineering leads with deep payments domain expertise (founders table). • 30-day web traffic up 32%, suggesting product-market fit acceleration (key_signals). • TechCrunch coverage highlights enterprise pilot with a Fortune 500 bank (news).* |
| **4** | Bullets are data-grounded and relevant, but draw from only 3 of 4 tables, or one bullet is slightly generic. | *All bullets are strong but none reference news articles (only 3 of 4 tables used).* |
| **3** | Bullets present but mostly restate data without synthesizing investment relevance. Or draw from only 2 tables. | *• Company raised $12M. • Team has 15 people. • Website updated recently. (Data restated, not synthesized.)* |

| | | |
|---|---|---|
| **2** | Fewer than 2 bullets, or bullets contain generic reasoning rather than company-specific reasoning. | *• This is an exciting company in a large market. • The founders have relevant backgrounds.* |
| **1** | Section missing, or bullets contain hallucinated reasoning or are completely generic. | *Section absent or states "this company could disrupt the $X trillion market" with no data support.* |

## 4.4 D4: Company Snapshot Completeness (8%)

A formatted reference table with required fields: Founded, HQ, Employees, Products, Customers, Total Funding, Last Round. Missing fields must display “Not found in table”. Must include “last updated” timestamp.

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | All 7 required fields present. Missing fields explicitly marked “Not found in table.” Last updated timestamp present and correct. Values match input data exactly. | *Founded: 2021 \| HQ: San Francisco, CA \| Employees: 42 \| Products: Real-time fraud detection API \| Customers: B2B fintech \| Total Funding: $15.2M \| Last Round: $12M Series A (Mar 2024) \| Last updated: 2024-11-15* |
| **4** | All fields present and correct. Minor formatting inconsistency (e.g., date format varies) or “last updated” uses a slightly different label. | *All fields correct but says “Updated: Nov 2024” instead of the exact observed_at timestamp.* |
| **3** | 6 of 7 fields present. Missing field is silently omitted instead of marked “Not found in table.” Or one field has a minor transcription error. | *Customers field is absent (was not in data) but no “Not found in table” placeholder. All other fields correct.* |
| **2** | Multiple fields missing or silently omitted. Or a field contains an incorrect value from the data. | *Only 4 of 7 fields present. No “Not found in table” markers. Last updated missing.* |
| **1** | Section missing, or contains hallucinated field values not in the input data. | *Lists a founding date that does not match any input source.* |

## 4.5 D5: Founder Information Quality (8%)

One entry per founder. Format: Name – Title | LinkedIn link + 2–3 sentence background. Missing backgrounds must say “Background not yet available.” If no founders: must say “No founder data available.”

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | All founders listed with correct names, titles, LinkedIn URLs. Each has a 2–3 sentence background from Swarm. Missing backgrounds explicitly flagged. Backgrounds are informative summaries and not raw data. | ***Jane Doe** – CEO & Co-Founder \| [LinkedIn](url)Former VP of Engineering at Stripe, where she led the payments infrastructure team. Stanford CS ’18. Previously built real-time systems at Google.* |
| **4** | All founders listed correctly. Backgrounds present but one is slightly thin (1 sentence instead of 2–3) or lacks specificity. | *Background says “Previously worked in fintech” when Swarm data specified “Former VP of Engineering at Stripe.”* |

| | | |
|---|---|---|
| **3** | Founders listed but one background is omitted (no “Background not yet available” marker), or a title is slightly wrong. | *Two founders listed but second founder’s background is missing with no acknowledgment.* |
| **2** | Significant omissions: a founder in the data is entirely missing from the output, or backgrounds contain errors. | *Data has 3 founders; output only lists 2 with no explanation.* |
| **1** | Section missing, says “No founder data available” when data exists, or contains fabricated founder backgrounds. | *Lists a founder credential (e.g., “PhD from MIT”) not present in any input source.* |

## 4.6 D6: Key Signals Quality (10%)

One bullet per signal from key_signals table. Website signals should be summarized concisely (not page titles repeated). Each signal must include source and timestamp. If no signals: must say “Source not yet implemented.”

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | All signals from the table are represented. Each is concisely described, includes source and timestamp. Website signals are summarized (not raw page titles). Investment relevance is clear. | *• Web traffic: +5.2% over 30 days, suggesting growing organic demand (Source: harmonic, 2024-11-15). • Website product: Added enterprise SSO and audit logging to pricing page, signaling upmarket move (Source: tavily, 2024-11-10).* |
| **4** | All signals present with sources. One signal’s description is slightly vague or redundant with the TL;DR. | *All signals listed but web_traffic signal just says “traffic is increasing” without the specific +5.2% figure.* |
| **3** | Most signals present. One signal is missing source/timestamp. Or a website signal repeats a raw page title instead of summarizing. | *• Website update: “/pricing page changed” (raw Tavily output, not synthesized into what changed).* |
| **2** | Fewer than half of available signals are represented. Or descriptions are so vague they provide no actionable information. | *Only 2 of 6 available signals listed. No sources or timestamps on any.* |
| **1** | Section missing, says “Source not yet implemented” when signals exist, or contains fabricated signals. | *Lists a “partnership with Microsoft” signal not present in key_signals table.* |

## 4.7 D7: News Synthesis Quality (10%)

One entry per article. Format: Headline | Outlet | Date, URL, then a 2–3 sentence synthesized takeaway. Takeaway must be synthesized from excerpts, not copied verbatim or made up. If no news: must say “No recent news available (source not yet implemented).”

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | All articles represented. Each has correct metadata (headline, outlet, date, URL). Takeaways are synthesized: they distill the excerpt into investment-relevant insights in the agent’s own words. | *Takeaway: “The article reports Acme signed its first Fortune 500 client (JPMorgan) for a fraud detection pilot, representing a significant enterprise validation milestone. The deal reportedly has a $500K ACV.”* |
| **4** | All articles present with correct metadata. Takeaways are accurate but one is slightly shallow (summarizes without framing). | *Takeaway accurately describes what happened but doesn’t frame why the investor should care.* |

| | | |
|---|---|---|
| **3** | Most articles present. One takeaway is a near-verbatim copy of the excerpt (raw excerpt dumping) rather than synthesis. Or an article about a similarly-named company is included without flagging ambiguity. Or synthesis is shallow and over-framed as “investor-focused”. | *Takeaway lifts 2+ consecutive sentences directly from the article excerpt.* |
| **2** | Multiple takeaways are raw excerpt dumps, or articles about the wrong company are included uncritically. | *3 of 5 takeaways are just pasted excerpts. No synthesis attempted.* |
| **1** | Section missing, contains fabricated news, or says “No news” when articles were provided. | *Takeaway describes events not mentioned in any provided article excerpt.* |

## 4.8 D8: Meeting Prep Actionability (10%)

2–3 suggested agenda items/questions + 1 recommended next step. Every item must be grounded in the input data and not generic VC advice.

| Score | Behavioral Anchor | Concrete Example |
|---|---|---|
| **5** | 2–3 questions that are specific to this company's data (referencing specific signals, founder backgrounds, or news). 1 concrete next step. | *• Ask about the JPMorgan pilot conversion timeline (from news). • Probe the -9.7% headcount decline: is this efficiency or distress? (from hiring signal). • Explore why web traffic spiked 32% — is it organic or paid? (from web_traffic signal). Next step: Request customer references from the enterprise pilot.* |
| **4** | Questions are data-grounded and company-specific, but one is slightly generic or the next step is vague. | *Questions reference real data but next step says "schedule a follow-up" without specificity.* |
| **3** | Questions are reasonable but only loosely connected to specific data points. Could apply to many similar companies. | *• Ask about product roadmap. • Discuss funding plans. (Reasonable but not grounded in THIS company's specific signals.)* |
| **2** | Questions are generic VC boilerplate with no company-specific grounding. | *• What's your TAM? • How do you think about competition? • What's your burn rate?* |
| **1** | Section missing, or contains questions based on hallucinated data, or is entirely generic. | *Questions reference a "competitor acquisition" not mentioned in any input data.* |

## 4.9 D9: Structural Compliance (5%)

Binary-style dimension: does the output follow the required 7-section structure, headings, and formatting rules?

| Score | Behavioral Anchor |
|---|---|
| **5** | All 7 sections present, in correct order, with correct headings. Markdown formatting is clean. Section lengths are within expected ranges. No structural anomalies. |
| **4** | All 7 sections present and in order. Minor formatting inconsistency (e.g., heading capitalization varies, extra blank line). |
| **3** | All 7 sections present but one is misordered or has a non-standard heading name. Or one section violates its length constraint (e.g., TL;DR exceeds 5 sentences). |
| **2** | One section is missing. Or two sections are misordered. Or output is truncated (under 500 characters). |

| | |
|---|---|
| **1** | Multiple sections missing. Or output is empty/near-empty. Or a fundamentally different structure was used (e.g., prose essay instead of the 7-section template). Automatic hard failure if ≥3 sections absent. |

## 4.10 D10: Clarity & Readability (5%)

Professional writing quality for a time-constrained reader. The investor should be able to scan the entire document in under 5 minutes.

| Score | Behavioral Anchor |
|---|---|
| **5** | Crisp and professional. Scannable: bullets are concise, paragraphs are short. Consistent formatting throughout. No grammatical errors. |
| **4** | Well-written with one minor issue: a slightly awkward sentence, an unclear abbreviation, or a formatting inconsistency. |
| **3** | Acceptable but not polished. Some sentences are overly long. Minor grammatical errors. Readability is adequate but not optimized for speed. |
| **2** | Difficult to scan quickly. Multiple long or unclear sentences. Inconsistent tone (mixes formal and informal). Multiple grammatical errors. |
| **1** | Incoherent writing, excessive jargon, or unreadable formatting. Or the document reads like raw LLM output (e.g., includes system prompt artifacts, placeholder text, or “As an AI language model” patterns). |

# 5. Hard Failure Conditions

The following conditions constitute automatic disqualification. If any hard failure is detected, the output receives a composite score of 0 regardless of performance on other dimensions. Record the specific failure type.

| ID | Failure | Definition & Detection |
|---|---|---|
| H F1 | Hallucination | Any factual claim in the output that cannot be traced to any field in the four input data tables. A single instance is sufficient. Detection: cross-reference every specific claim (names, numbers, dates, events) against input data. |
| H F2 | Outside Knowledge | The agent uses information from its training data rather than the input tables. Indicators: mentions of events, competitors, or market data not present in any input source. |
| H F3 | Wrong Company | The briefing describes a different company than the one requested. Maps to the wrong_company error type and Entity Resolution = Incorrect. |
| H F4 | Empty/Truncated Output | The agent returned no content (empty error type) or fewer than 500 characters (truncated). Maps directly to the automated LLM Error Detection check. |
| H F5 | Forbidden Pattern | Output contains LLM refusal patterns (“I cannot help with that”), system prompt leakage, or apology patterns that indicate a generation failure. |

# 6. Failure Mode Severity Mapping

| Error Type | Severity | Dimension | Scoring Impact |
|---|---|---|---|
| **empty** | Critical | D9 | Hard failure (HF4). Composite = 0. |
| **truncated** | Critical | D9 | Hard failure (HF4). Composite = 0. |
| **missing_sections** | Major | D9 | D9 = 1 if ≥3 sections absent (hard failure). D9 = 2 if 1–2 sections absent. |
| **hallucination** | Critical | D1 | Hard failure (HF1). D1 = 1 automatically. |
| **outside_knowledge** | Critical | D1 | Hard failure (HF2). D1 = 1 automatically. |
| **echo** | Major | D10 | Deduct 2 points from the affected section dimension + D10. |
| **wrong_company** | Critical | D1 | Hard failure (HF3). Composite = 0. |
| **stale_data** | Minor | D1 | Note in context tags. Do not penalize unless agent misrepresents recency. |

| | | | |
|---|---|---|---|
| **missing_background** | **Moderate** | D5 | D5 deducted 1–2 points depending on number of silently omitted backgrounds. |
| **raw_excerpt** | **Major** | D7 | D7 deducted 2 points per raw excerpt dump. |
| **generic_signals** | **Moderate** | D6 | D6 deducted 1–2 points based on vagueness. |
| **generic_meeting_prep** | **Major** | D8 | D8 ≤ 2 if all questions are generic boilerplate. |
| **missing_data_silent** | **Moderate** | D4/D5 | Deduct 1 point from affected section dimension per silently omitted field. |
| **json_parse** | **Moderate** | D9 | D9 deducted 1–2 points based on severity of formatting error. |
| **forbidden_pattern** | **Critical** | D9 | Hard failure (HF5). Composite = 0. |

# 7. Edge Case Handling

When data is legitimately missing vs. when the agent fails to surface available data, the scoring treatment is different. This section provides rules.

| Scenario | If Data Was Unavailable | If Data Was Available But Missed |
|---|---|---|
| **No news articles returned by Parallel Search** | Agent says “No recent news available (source not yet implemented).” Score D7 = 5 (correct behavior). | Agent says “No recent news” but articles existed in input. Score D7 = 1. |
| **No founder data in Harmonic or Swarm** | Agent says “No founder data available.” Score D5 = 5. | Agent omits founder section or says “No data” when founders were in the input. Score D5 = 1. |
| **Harmonic returns no products field** | Agent marks “Not found in table” in Snapshot. TL;DR Sentence 1 uses whatever alternative description is available. Score D4 = 5. | Agent silently omits Products from Snapshot. Score D4 ≤ 3 (missing_data_silent penalty). |
| **Only 1 signal type in key_signals** | Agent renders the single signal with source and timestamp. Score D6 relative to available data (4–5 if well-rendered). | Agent has 6 signals in data but only renders 1. Score D6 = 2. |
| **Tavily returns 0 results (site blocks crawlers)** | Record Signal Coverage = Partial or Minimal in context tags. Website signals absent — no penalty on D6 for missing website-sourced signals. | N/A — evaluator cannot distinguish (data unavailability is upstream). |
| **Article is about a similarly-named company** | Agent includes it without flagging ambiguity. Score D7 deducted 1 point for failing to caveat. Score D1 not penalized (Parallel Search retrieval issue, not hallucination). | Agent notices and flags “Note: this article may reference a different entity.” Score D7 = 5. |
| **Company has no funding history** | Total Funding and Last Round fields show “Not found in table.” Score D4 = 5. | Agent invents a funding round. Hard failure (HF1). Score D1 = 1. |

The agent should never be penalized for data that the upstream pipeline didn’t provide. It should always be penalized for failing to use data that was provided, or for silently omitting fields that should be marked “Not found in table.”

# LLM-as-a-Judge Implementation Notes

### Prompt Architecture

- **Inputs to the judge:** (1) the agent's output, (2) the complete input data tables (all four sources), and (3) this rubric's dimension definitions with Likert anchors.
- **Chain-of-thought:** Require the judge to output reasoning per dimension before scoring. This follows the G-Eval pattern and significantly improves scoring reliability.
- **Structured JSON output:** { dimensions: { D1: { score, reasoning, violations }, ... }, composite_score, hard_failures: [], context_tags: {} }
- **Hard failure first:** Instruct the judge to scan for HF1–HF5 before scoring dimensions. If any hard failure is detected, return composite = 0 immediately.

---

# Relationship to Existing Metrics in our Agent

The system already computes five automated metrics. This rubric does not duplicate them but uses them as inputs to the evaluation workflow.

| Automated Metric | Rubric Relationship | How to Use |
|---|---|---|
| **Entity Resolution Accuracy (0–1)** | Feeds context tag: Entity Resolution | If < 0.8, tag as Ambiguous. If wrong_company detected, output is a hard failure (HF3) before rubric scoring begins. |
| **Signal Coverage Rate (0–1)** | Feeds context tag: Signal Coverage | Record as Full/Partial/Minimal. Adjusts expectations for D6 (Key Signals Quality). Low coverage = lower ceiling for D6. |
| **Retrieval Precision (per source)** | Informs D1 and D7 scoring | Low Parallel Search precision means some articles may be about wrong companies. Evaluator should check before penalizing D7. |
| **Structural Validation (pass/fail)** | Pre-screen for D9 | If structural validation = fail, D9 starts at ≤2. Evaluator still scores other dimensions independently. |
| **LLM Error Detection** | Pre-screen for hard failures | If empty, truncated, or forbidden_pattern detected, output is a hard failure (HF4 or HF5) before rubric scoring begins. |