NEWS AGGREGATOR AGENT

Scoring Rubric

Prepared by MS&E108 Group

1. Purpose and Scope

Unlike a pure LLM agent, most decisions here are algorithmic (scoring, filtering, deduplication, classification).

There are three important aspects we want to highlight:

Quality bar: Signal-to-noise ratio. Every signal in the digest must be there for a reason. False positives (irrelevant signals surfaced) and false negatives (important signals missed) are the two failure modes of our agent.

Two-tier structure: The rubric is split into Tier A (is each signal correctly fetched, classified, scored, and described?) and Tier B (is the digest well-curated, well-ordered, and investor-ready?). A signal can be individually pretty good while a digest is still bad if selection/ordering logic fails, and the other way around.

Separation of concerns: We have three independent axes. Pipeline Accuracy (did filtering, classification, deduplication, scoring work correctly). Output Quality (how well the digest was assembled and synopses written). Data Availability (how much signal the upstream APIs returned). A failure shouldn't be penalized on more than one axis.

---

2. Evaluation Context Tags

We (or LLM judge) record these before scoring. They normalize comparisons across different run conditions.

Tag	Values	Why It Matters
Run Mode	digest / alerts / signals	Alert mode evaluates only FUNDING, M&A, team_change. Digest mode evaluates the full pipeline. Signals mode has no minimum score filter.
Lookback Window	N days (default 7)	Shorter windows reduce available signals and affect recall expectations.
API Coverage	Full / Partial / Degraded	Full = all APIs returned data. Partial = 1+ APIs timed out or returned empty. Degraded = majority of calls failed. Record which APIs were affected.
Watchlist Size	N companies (portfolio + competitor)	Larger watchlists increase the difficulty of curation and the risk of misattribution.

LLM Mode	Enabled / Disabled (--fast)	When LLM is disabled, all synopses are template-based. Synopsis hallucination checks are N/A.
Min Priority Threshold	Score (default 40)	Lower thresholds let in more noise; higher thresholds can risk missing important signals.
Embedding Availability	Available / Unavailable	When embeddings are unavailable, classification falls back to hash-based pseudo-embeddings and is unreliable. Score classification dimension with more leniency.

3. Global Scoring Scale

Each dimension is scored 1–5. The composite is a weighted sum normalized to 0–100.

Score	Label	Definition
5	**Excellent**	Every featured signal is relevant, correctly classified, properly sourced. No false positives in the digest view. Major events for all portfolio companies are captured. Synopsis quality is good and investment-relevant.
4	**Good**	Usable with minor issues. 1–2 borderline signals included or one notable signal buried below the fold. All critical events captured. Minor classification or sentiment error on a non-featured signal.
3	**Acceptable**	Covers the important events but has noticeable quality gaps: a duplicate not merged, a competitor signal outranking a similar portfolio signal, or a synopsis that's too vague. Investor would need to cross-reference.
2	**Below Standard**	Significant problems: multiple misattributed signals, a major event missing from featured stories, homepage URL in the digest, or widespread classification errors. Not trustworthy without manual review.
1	**Poor**	Fundamental failure: wrong company signals dominate the digest, portfolio company with major news has zero signals, hallucinated synopsis content, or digest is empty/structurally broken.

4. Signal-Type × Company-Category Severity Matrix

Not all missed or mishandled signals are equally severe. This matrix encodes the principle that a missed funding signal for a portfolio company is a lot worse than a missed market signal for a competitor. Use this to calibrate deductions across all dimensions when you score.

Signal Type	Portfolio Company	Competitor

FUNDING / M&A / IPO	CRITICAL	HIGH
SECURITY / LEGAL / LAYOFFS	CRITICAL	HIGH
HIRING / team_change (C-suite)	HIGH	MEDIUM
PRODUCT / PARTNERSHIP / CUSTOMER	MEDIUM	MEDIUM
EARNINGS / MARKET / GENERAL	LOW	LOW

5. Scoring Dimensions

12 dimensions across two tiers. Tier A evaluates individual signal quality (pipeline accuracy). Tier B evaluates digest assembly (output quality). Weights reflect the investor use case. Precision/recall carry the most weight, then classification and deduplication, then formatting.

ID	Tier	Dimension	Wt	Axis	Rationale
S1	A	**Signal Attribution Accuracy**	15%	Pipeline	A signal attributed to the wrong company is the most damaging error and obviously misleads investor decisions.
S2	A	**Filter Quality (Noise Rejection)**	10%	Pipeline	Homepages, press releases, and listicles in the digest waste time and erode trust
S3	A	**Classification Accuracy**	10%	Pipeline	A FUNDING article classified as PRODUCT distorts the priority ordering and interpretation.
S4	A	**Deduplication Quality**	8%	Pipeline	Same event appearing twice wastes time. Merging distinct events is also harmful.
S5	A	**Scoring & Ranking Accuracy**	8%	Pipeline	Incorrect ranking buries important signals below less important ones.
S6	A	**Sentiment Accuracy**	4%	Pipeline	Wrong sentiment misleads investor interpretation but is lower-stakes than attribution or classification.
S7	A	**Synopsis Quality (LLM)**	8%	Output	LLM-enhanced synopses have a hallucination risk
S8	A	**Synopsis Quality (Template)**	4%	Output	Template synopses can be too vague. Lower weight because hallucination risk is near-zero.
D1	B	**Digest Precision**	12%	Output	% of signals in digest that are truly relevant.
D2	B	**Digest Recall**	10%	Output	% of major real-world events surfaced. Evaluated against the severity matrix!
D3	B	**Digest Ordering & Curation**	6%	Output	Are the right stories featured? Is portfolio prioritized? Is recency reflected?
D4	B	**Structural & Format Compliance**	5%	Output	Correct Markdown structure, source links present, dates present, section limits respected.

Composite: Score = [(S1×15)+(S2×10)+(S3×10)+(S4×8)+(S5×8)+(S6×4)+(S7×8)+(S8×4)+(D1×12)+(D2×10)+(D3×6)+(D4×5)] / 5, yielding 0–100.

5.1 TIER A: Individual Signal Quality

S1: Signal Attribution Accuracy (15%)

Is each signal correctly attributed to the right company? This is the single most critical pipeline dimension. A wrong-company signal is worse than a missing signal.

Score	Behavioral Anchor	Concrete Example
5	100% of signals in the digest are correctly attributed to their target company. No cross-industry name collisions leaked. Known ambiguous names (Sana, Merge) handled correctly.	*Digest for "Sana Labs" contains zero articles about Sana Biotechnology.*
4	99%+ attribution accuracy. One borderline signal where the article mentions the target company only tangentially (appears in a competitor list, not as the subject).	*Article titled "Top 10 AI startups" includes the target company in a list; signal is attributed but article isn't really "about" the company.*
3	95%+ accuracy. 1–2 signals attributed to the wrong company, but the errors are in non-featured positions and involve genuinely ambiguous names.	*One signal about "Merge Labs" (BCI company) leaks into the digest for "Merge" (API platform) in the More Headlines section.*
2	90–94% accuracy. Multiple misattributed signals, or one misattributed signal in the Featured Stories section.	*A featured story about "Stripe" is actually about "Stripe Climate" (a different product vertical), leading to a misleading synopsis.*
1	Below 90% accuracy, or a Featured Story is about an entirely wrong company. Hard failure.	*The #1 featured story for a portfolio company is actually about a competitor with a similar name.*

Hard failure: Any Featured Story attributed to the wrong company = automatic S1 = 1 and composite = 0.

S2: Filter Quality / Noise Rejection (10%)

Did the filter stage correctly remove homepages, press releases, Wikipedia, listicles, stale articles, and non-article URLs? Any of these appearing in the digest is a filter failure.

Score	Behavioral Anchor	Concrete Example
5	Zero noise signals in the digest. All homepage URLs, press releases, Wikipedia articles, listicles, and stale articles were correctly filtered.	*No signals with URLs matching /category/, /tag/, /press-releases/, or known non-article patterns.*

4	Zero noise in Featured Stories or More Headlines. One borderline item (e.g., a company blog post that is half press release, half genuine news) in industry trends.	*A Crunchbase profile page appears at position 11 in a long signal list but not in the digest view.*
3	One noise signal leaked into More Headlines (not Featured). Or one stale article (>7 days) appears in the digest.	*A "Press Release: Company X Announces..." article appears at position 6 in More Headlines.*
2	Multiple noise signals in the digest, or one noise signal in Featured Stories.	*A Wikipedia article about the company appears as Featured Story #3.*
1	A lot of noise: 3+ non-article URLs, press releases, or stale signals in the digest. Hard failure if homepage URL appears in Featured.	*Homepage URL (stripe.com with no path) appears as a featured story.*

S3: Classification Accuracy (10%)

Is each signal assigned the correct signal type?

Score	Behavioral Anchor	Concrete Example
5	All signals in the digest have the correct signal type. Rule-based classifications match the article content. Embedding-based classifications are appropriate.	*"Company raises $50M Series B" correctly classified as FUNDING with confidence 0.95.*
4	One minor misclassification on a non-featured signal where the distinction is ambiguous (e.g., PARTNERSHIP vs. CUSTOMER for a joint venture announcement).	*A "strategic partnership" that includes a customer deal component is classified as PARTNERSHIP instead of CUSTOMER. Reasonable ambiguity.*
3	One misclassification on a featured signal (e.g., a CEO hire classified as GENERAL instead of HIRING), or 2–3 misclassifications on non-featured signals.	*CEO appointment article classified as GENERAL because rule patterns didn't match the phrasing "elevated to CEO."*
2	Multiple featured signals misclassified, or a CRITICAL-severity signal type is wrong (FUNDING classified as PRODUCT).	*"Company acquired by BigCorp for $2B" classified as PARTNERSHIP instead of M&A.*
1	Widespread misclassification (>20% of digest signals wrong). Or a FUNDING/M&A signal classified as GENERAL and consequently scored below the digest threshold, making it invisible. This is a hard failure.	*A $100M Series C for a portfolio company is classified GENERAL, scored at 35, and excluded from the digest entirely.*

S4: Deduplication Quality (8%)

Are same-event articles correctly merged? Are distinct events incorrectly merged? Source selection for primary URL matters.

Score	Behavioral Anchor	Concrete Example
5	All same-event articles merged. No false merges (distinct events incorrectly combined). Primary URL selected from highest-quality source.	*TechCrunch, Bloomberg, and Reuters articles about the same $50M raise are merged. TechCrunch selected as primary (highest priority score).*
4	All merges correct. One primary URL selection is not optimal (a blog post selected over TechCrunch due to URL normalization artifact) but the correct article is still accessible.	*Primary URL is from a lesser-known outlet when TechCrunch also covered the story, due to URL normalization selecting the blog's cleaner URL.*

3	One duplicate not merged (same event appears twice in the digest), or one false merge (two distinct events combined into one signal).	*Reuters and Bloomberg articles about the same funding round appear as separate signals in More Headlines.*
2	Multiple duplicates visible in the digest. Or a false merge combined a funding announcement with an unrelated product launch because title similarity exceeded 60%.	*Two featured stories are actually the same M&A announcement from different outlets.*
1	Deduplication essentially non-functional: 3+ duplicate events in Featured Stories, or critical false merges that erase distinct signals.	*Featured Stories #1 and #2 are the same event. A funding announcement was merged with a layoff article because both headlines mentioned the company name.*

S5: Scoring & Ranking Accuracy (8%)

Are signals ranked in the correct order?

Score	Behavioral Anchor	Concrete Example
5	Featured stories are the 3 most important signals. Portfolio companies correctly prioritized. Recency bonus properly applied. Score breakdowns are consistent with the documented formula.	*Portfolio company's $50M funding (rank=95) correctly featured above competitor's product launch (rank=65).*
4	Ranking is correct for Featured Stories. One minor ordering issue in More Headlines (e.g., a 5-day-old story ranked above a 1-day-old story of similar importance).	*More Headlines has a minor recency inversion between two MEDIUM-severity signals.*
3	One Featured Story is arguably not top-3. Or a competitor signal of medium importance outranks a portfolio signal of equal importance.	*A competitor's product launch is Featured Story #3 when a portfolio company's partnership announcement (same day, similar importance) is in More Headlines.*
2	A clearly low-importance signal (MARKET or GENERAL) appears in Featured Stories while a CRITICAL-severity signal is buried in More Headlines.	*GENERAL industry overview article featured while a portfolio company's Series C is at position 7.*
1	Ranking is random. Critical portfolio signals missing from top positions. Score inflation.	*A competitor's listicle mention (rank=85 due to scoring bug) outranks a portfolio company's acquisition announcement.*

S6: Sentiment Accuracy (4%)

Is the sentiment label (positive/negative/neutral) correct for each signal? Lower weight because sentiment is supplementary information and not a primary decision driver.

Score	Behavioral Anchor
5	All signals have correct sentiment. Funding correctly marked positive. Layoffs correctly marked negative. Ambiguous events correctly marked neutral.
4	One minor sentiment error on a genuinely ambiguous signal (e.g., "acquisition" marked neutral when it could be positive or negative depending on context).
3	1–2 sentiment errors on non-featured signals. Or a featured signal's sentiment is correct but the wrong intensity (a massive funding round marked neutral instead of positive).
2	A featured signal has inverted sentiment (layoff article marked positive, or funding article marked negative).

1	Widespread sentiment errors (>30% wrong). Or a portfolio company's critical negative event (data breach) is marked positive.

S7: Synopsis Quality — LLM-Enhanced (8%)

Applies only when LLM mode is enabled and synopsis was LLM-generated (featured stories with confidence ≥ 0.7). If LLM is disabled, score N/A and redistribute weight equally to S8 and D1.

Score	Behavioral Anchor	Concrete Example
5	Synopsis is 1–2 sentences. States what happened, why it matters, and includes key numbers. Every claim is traceable to the headline/excerpt. Investment-relevant framing and concise.	*"Stripe raised a $50M Series B led by Sequoia, bringing total funding to $75M. The round suggests strong enterprise traction following a 3x revenue year." (Numbers match excerpt.)*
4	Synopsis is accurate and concise. Slightly generic framing ("this is significant") but no hallucination. All facts traceable.	*Synopsis accurately describes the event but says "this could be meaningful for the space" instead of citing a specific metric from the excerpt.*
3	Synopsis is accurate but reads as a paraphrase of the headline (low synthesis effort). Or is 3–4 sentences (slightly too long).	*Synopsis essentially restates the headline with minor wording changes and no investment framing.*
2	Synopsis contains one claim not clearly supported by the headline/excerpt (borderline hallucination). Or is an echo of the input.	*Synopsis mentions "a $2B valuation" when the excerpt only says "raised $50M" with no valuation mentioned.*
1	Synopsis contains a fabricated fact (clear hallucination). Automatic hard failure.	*Synopsis claims "the acquisition was completed for $500M" when the article only mentions "acquisition talks."*

Any hallucinated fact in an LLM synopsis = S7 = 1 and automatic flag. If the hallucinated synopsis appears in a Featured Story, composite = 0.

S8: Synopsis Quality — Template-Based (4%)

Template synopses have near-zero hallucination risk but can be too vague or fail to extract key details.

Score	Behavioral Anchor	Concrete Example
5	Template correctly populated all slots: company name, amount, series, context. Synopsis is specific and informative despite being formulaic.	*"Stripe raised $50M in Series B funding. The round was led by Sequoia Capital." (All slots filled from article data.)*

4	Template mostly populated. One optional slot missing (e.g., lead investor not extracted) but core information present.	*“Stripe raised $50M in Series B funding.” (Lead investor not extracted but available in excerpt.)*
3	Template produced a vague result because extraction failed on key details. Synopsis is technically correct but uninformative.	*“Stripe announced funding details.” (Amount and series not extracted.)*
2	Template produced a near-empty synopsis (“Company announced details.”) for a signal where rich data was available in the excerpt.	*“Acme announced details.” for a $100M Series C with extensive excerpt data.*
1	Template synopsis is factually wrong (wrong amount extracted, wrong company name inserted) or completely empty.	*Template says “raised $5M” when excerpt clearly states “$50M.” (Regex extraction error.)*

5.2 TIER B: Digest Assembly Quality

D1: Digest Precision (12%)

What percentage of signals in the digest view are truly relevant to the target companies? Evaluated on the full digest (Featured + More Headlines + Industry Trends).

Score	Behavioral Anchor	Concrete Example
5	Precision ≥ 95%. Every signal in the digest is clearly about the correct company, is a genuine news article (not noise), and is within the date range.	*20 signals in digest; 20/20 are relevant, correctly attributed, and timely.*
4	Precision 90–94%. 1–2 borderline signals (tangential mention, or company appears only in a list within the article).	*20 signals; 18–19 clearly relevant; 1–2 are "about" the company but only tangentially.*
3	Precision 80–89%. 3–4 irrelevant or borderline signals. Noise is visible but concentrated in More Headlines, not Featured.	*20 signals; 4 are press releases or tangential articles that passed the filter.*
2	Precision 70–79%. Noise is visible in Featured Stories or comprises >20% of digest.	*15 signals; 4 are about wrong companies or are non-article URLs.*
1	Precision < 70%. The digest is more noise than signal. Investor cannot trust it. Hard failure.	*10 signals; 4 are about wrong companies, 2 are Wikipedia articles.*

D2: Digest Recall (10%)

What percentage of major real-world events surfaced? Evaluated against the severity matrix: critical events for portfolio companies must be captured. Use external verification (manual check of company news) to assess recall!

Score	Behavioral Anchor	Concrete Example
5	Recall ≥ 90%. All critical events for portfolio companies captured. All high events for portfolio and competitors captured. At most 1 medium event missed.	*Portfolio company announced Series C (captured), competitor had layoffs (captured), portfolio company hired new CTO (captured).*
4	Recall 80–89%. All critical events captured. 1–2 high events missed (e.g., a competitor's VP hire not surfaced).	*All funding and M&A events captured; one competitor's product launch missed because Parallel Search didn't return it.*

3	Recall 70–79%. One critical event for a competitor missed, or 2–3 high events missed. All critical portfolio events captured.	*Competitor's $200M acquisition missed entirely. But all portfolio company events captured.*
2	Recall 60–69%. One critical event for a portfolio company missed (but the API returned the article and it was filtered incorrectly).	*Portfolio company's Series B was in the raw fetch but filtered out because the company name appeared only in the excerpt, not the headline.*
1	Recall < 60%, or a critical portfolio event was missed entirely despite being available in the API data. Hard failure.	*Portfolio company announced acquisition; Parallel Search returned the TechCrunch article; filter excluded it. Investor learned about it from Twitter instead.*

On data availability: If Parallel Search did not return any articles about an event (API Coverage = Partial/Degraded), the missed event is a data availability issue, not a digest quality failure. Record in context tags and do not penalize D2.

D3: Digest Ordering & Curation (6%)

Are Featured Stories the right 3? Is portfolio prioritized over competitors of equal importance? Is recency reflected?

Score	Behavioral Anchor
5	Featured Stories are unambiguously the 3 most important signals. Portfolio signals ranked above competitor signals of equal importance. Recency properly reflected. Industry Trends section adds useful context. Overview stats are also accurate.
4	Featured Stories are correct. One minor curation issue: a competitor signal ranked above a portfolio signal of similar (not identical) importance, or Industry Trends section is thin.
3	One Featured Story is arguably not top-3 (a more important signal is in More Headlines). Or portfolio/competitor prioritization is inconsistent. Or overview stats are slightly off.
2	Featured Stories include a clearly low-priority signal while a CRITICAL event is buried. Or the digest is competitor-heavy when portfolio companies had significant news.
1	Featured Stories are essentially random. No coherent prioritization. Portfolio company with major news has zero featured signals while competitors dominate.

D4: Structural & Format Compliance (5%)

Does the digest follow the documented Markdown format? Are source links present? Dates present? Section limits respected?

Score	Behavioral Anchor
5	All sections present (Overview, Featured Stories, More Headlines, optionally Industry Trends). Source URLs present and functional for every signal. Published dates present. Section caps respected (≤3 featured, ≤8 more headlines, ≤5 industries). Generation timestamp present.
4	All sections present. One minor formatting issue (e.g., sentiment icon missing, date format inconsistent). All source URLs present.
3	One non-critical section missing (e.g., Industry Trends). Or 1–2 signals missing source URLs. Or section caps slightly exceeded.
2	Multiple format violations: missing source URLs for featured stories, no overview stats, or significantly exceeded section caps.
1	Digest is structurally broken: missing Featured Stories section, empty output, or output in wrong format (JSON instead of Markdown for --digest mode).

6. Hard Failure Conditions

If any hard failure is detected, the output receives composite = 0 regardless of other dimensions.

ID	Failure	Definition
HF1	Synopsis Hallucination	An LLM-generated synopsis contains a factual claim (number, date, company name, event) not present in the article headline or excerpt. A single instance in a Featured Story is sufficient.
HF2	Featured Wrong-Company	A Featured Story (top 3) is attributed to the wrong company. The article is clearly about a different entity than the target company.
HF3	Portfolio CRITICAL Miss	A CRITICAL-severity event (FUNDING, M&A, IPO, SECURITY, LEGAL, LAYOFFS) for a portfolio company was present in the raw API data but absent from the digest due to a pipeline failure (filter, classification, or scoring error).
HF4	Homepage in Digest	A homepage URL (domain with no article path) appears anywhere in the digest view (Featured or More Headlines).
HF5	Empty Digest	The digest is empty, truncated, or structurally broken (no parseable sections) when the raw data contained signals above the minimum threshold.

7. Failure Mode Severity Mapping

Error Type	Severity	Dim	Impact
wrong_company	**Critical**	S1	S1 = 1. If in Featured, HF2 → composite = 0.
homepage_leaked	**High**	S2	S2 deduct 2. If in Featured, HF4 → composite = 0.
noise_leaked	**High**	S2	S2 deduct 1–2 per noise signal in digest.
stale_signal	**High**	S2	S2 deduct 1 per stale signal. Featured stale = deduct 2.
duplicate_not_merged	**High**	S4	S4 deduct 1 per unmerged duplicate. Featured duplicate = deduct 2.
wrong_classification	**High**	S3	S3 deduct 1–2 per misclassification. CRITICAL type misclassified = deduct 2.
low_confidence_featured	**Medium**	S3	Note in scoring. Do not penalize if rule-based classification was correct.
sentiment_mismatch	**Medium**	S6	S6 deduct 1 per mismatch. Inverted sentiment on Featured = deduct 2.

score_inflation	**Medium**	S5	S5 deduct 1–2 per inflation that affected Featured selection.
competitor_over_portfolio	**Medium**	S5/D3	S5 deduct 1. D3 deduct 1 if it displaced a portfolio Featured story.
recency_not_reflected	**Low**	S5	S5 deduct 0.5 per inversion in same-importance tier.
synopsis_hallucination	**Critical**	S7	S7 = 1. If in Featured, HF1 → composite = 0.
synopsis_echo	**High**	S7	S7 deduct 2. Verbatim echo is a synthesis failure.
synopsis_generic	**Medium**	S7/S8	S7 or S8 deduct 1. Vague but accurate is better than hallucinated.
synopsis_too_long	**Low**	S7/S8	Deduct 0.5. Minor formatting concern.
missing_source_link	**High**	D4	D4 deduct 1 per missing link. Featured missing link = deduct 2.
missing_date	**Medium**	D4	D4 deduct 0.5 per missing date.
industry_misattributed	**Medium**	D3	D3 deduct 1 per misattributed industry tag.
competitor_not_discovered	**High**	D2	D2 deduct 1–2 depending on competitor's significance.
no_signals_generated	**Critical**	D2	If portfolio company: HF3 if events existed. D2 = 1.
template_synopsis_generic	**Low**	S8	S8 deduct 0.5. Expected limitation of template approach.
date_parse_failure	**Medium**	S2	S2 deduct 1 if it caused a valid article to be excluded.

9. Edge Case Decision Tree

Specific edge cases documented in the system with scoring guidance.

Scenario	Correct Behavior	Scoring Guidance
Company has no news in the time window (API returned 0 articles)	Digest omits the company or notes “No recent signals.” This is correct behavior when data is unavailable.	Do not penalize D2. Record API Coverage = Partial. If ALL companies have 0 signals, investigate API failure.
Cross-industry name collision: filter correctly excludes “Sana Biotechnology” article when tracking “Sana Labs”	Article is filtered out. No signal generated. This is correct.	S1 = 5 and S2 = 5 for this event. Do not penalize recall — the article was correctly identified as irrelevant.
Cross-industry name collision: filter FAILS and includes wrong-company article	Wrong-company signal appears in digest.	S1 deducted. If in Featured, HF2 triggered.
Competitor signals appear but parent portfolio company has 0 signals	This can be correct (portfolio company had no news) or a failure (portfolio company had news that was missed).	Check raw API data. If portfolio company had articles in fetch: penalize D2 and S2 (filter failure). If API returned nothing: record in context tags, no penalty.
Featured story has confidence = 0.65 (below 0.7 LLM threshold), classified correctly by rules	Signal gets template synopsis instead of LLM synopsis. Classification is correct. This is expected behavior.	Do not penalize S3 or S7. Score S8 based on template quality. Note: confidence < 0.7 is a classification characteristic, not a failure.
Deduplicator selects lower-quality source as primary (blog post over TechCrunch) due to URL normalization	The correct event is represented but with a suboptimal source link.	S4 = 4 (merge was correct, source selection was suboptimal). D4 not penalized unless source link is broken.
A competitor’s FUNDING announcement is the #1 story when a portfolio company’s PRODUCT launch also occurred	FUNDING (+40 importance) should outrank PRODUCT (+20) regardless of company category, because signal importance > company type in the scoring formula.	This is CORRECT behavior per the scoring formula. S5 = 5. Do not penalize for competitor outranking portfolio when the signal type importance differential is large.
--min-priority set to 60 by investor, causing a portfolio PARTNERSHIP signal (rank=55) to be excluded	Signal correctly excluded per investor’s threshold configuration.	Record min-priority threshold in context tags. Do not penalize D2 for signals excluded by explicit investor configuration.

10. Scoring Protocol

Reference Materials Required

- **For all dimensions:** The digest output, the raw signal list (all fetched signals before filtering), and the scoring breakdown JSON.
- **For S1 (attribution) and D1 (precision):** The article headlines and excerpts. Evaluator must verify that each signal's article is actually about the claimed company.
- **For S7 (LLM synopsis):** The original article headline and excerpt. Evaluator must cross-reference every claim in the synopsis against the source material.
- **For D2 (recall):** External verification. A manual check of major company news for the time window (e.g., Crunchbase, TechCrunch, Google News). This is the only dimension that requires reference outside the system's own data and where a human is needed (talk to team about this).